

# **Concurrency Control**

**(Ref: Chapters 17 and 18,  
Silberchatz' Text)**

# INTRODUCTION

- In a Multiprogramming Environment
  - Multiple **transactions** (trans) may be executed concurrently
  - Need a **concurrency control scheme** to prevent transactions from destroying database consistency

# INTRODUCTION (2)

- Example: A “lost update” problem caused by uncontrolled concurrent transactions
  - Transaction  $T_1$ : Cancels  $N$  reservations from a flight (1) whose number of reserved seats is in  $X$  (Database); then reserves the same number of seats on another flight (2) whose number of reserved seats is in  $Y$ .
  - Transaction  $T_2$ : Reserves  $M$  seats on flight (1)

# INTRODUCTION (3)

- Example (cont.):

# SCHEDULES

- A **schedule**: a chronological order in which transactions are executed
- A schedule for a set of transactions:
  - Must consist of all instructions (operations) of those transactions
  - Must preserve the order of instructions within each transaction
- Example: in Transaction  $T_1$ 
  - if READ (X) appears before WRITE (Y)
    - => Any valid schedule that consists of  $T_1$ : READ (X) is performed **before** WRITE (Y) by  $T_1$ .

# SCHEDULES (2)

e.g.      **A=\$1000**      **B=\$2000**

**Trans T<sub>0</sub>: Transfer \$50 from A to B**

**T<sub>0</sub>: Read (A)**  
**A=A-50**  
**Write (A)**  
**Read (B)**  
**B=B+50**  
**Write (B)**

**Trans T<sub>1</sub>: Transfer 10% of balance from A to B**

**T<sub>1</sub>: Read (A)**  
**Temp:= A\*0.1**  
**A:=A-Temp**  
**Write (A)**  
**Read (B)**  
**B=B+ Temp**  
**Write (B)**

**Consistency Constraint: Total balance in A+B must be unchanged after execution of T<sub>0</sub>, T<sub>1</sub>.**

# SCHEDULES (3)

- **Questions:** what are the possible serial schedules? Are they correct schedules?
- **Answer:**

# SCHEDULES (4)

- **Question:** show one correct non-serial schedule.
- **Answer:**



# SCHEDULES (5)

- **Question:** show one incorrect non-serial schedule.
- **Answer:**

# SCHEDULES (6)

- Assumptions:
  - Every transaction is correct if executed on its own
  - Transactions do not depend on one another
  - Every **Serial Schedule** is correct
- REQUIREMENT: for concurrent execution
  - A schedule, after execution, must leave the database in a consistent state, i.e. it must be equivalent to some serial schedule.
  - i.e. it must be a **Serializable Schedule**

# SCHEDULES (7)

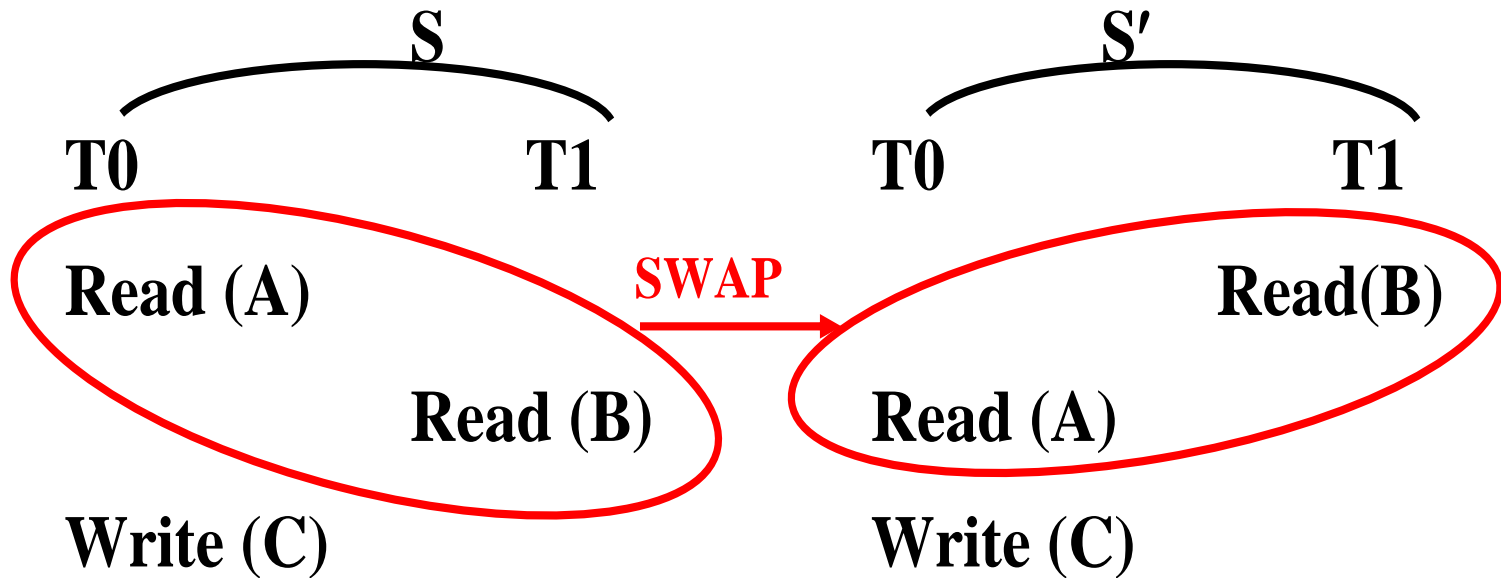
- Definition: A serializable schedule of n concurrent transactions: equivalent to some serial schedule of the same n transactions
- Otherwise => Non-serializable Schedule

# SCHEDULES (8)

- From a scheduling point of view, only significant operations of a trans are **READ** and **WRITE**
- Let  $I_i$  be an operation of Transaction  $T_i$
- Let  $I_j$  be an operation of Transaction  $T_j$
- $I_i$  and  $I_j$  are conflict if  $I_i$  and  $I_j$  access the same data item and either  $I_i$  or  $I_j$  is a WRITE operation (Read vs. Write, Write vs. Write)
- In a schedule  $S$ , if  $I_i$  and  $I_j$  are not conflict, then we can swap the order of  $I_i$  and  $I_j$  to produce a new schedule  $S'$ :  
 $S \equiv S'$  (Equivalent)

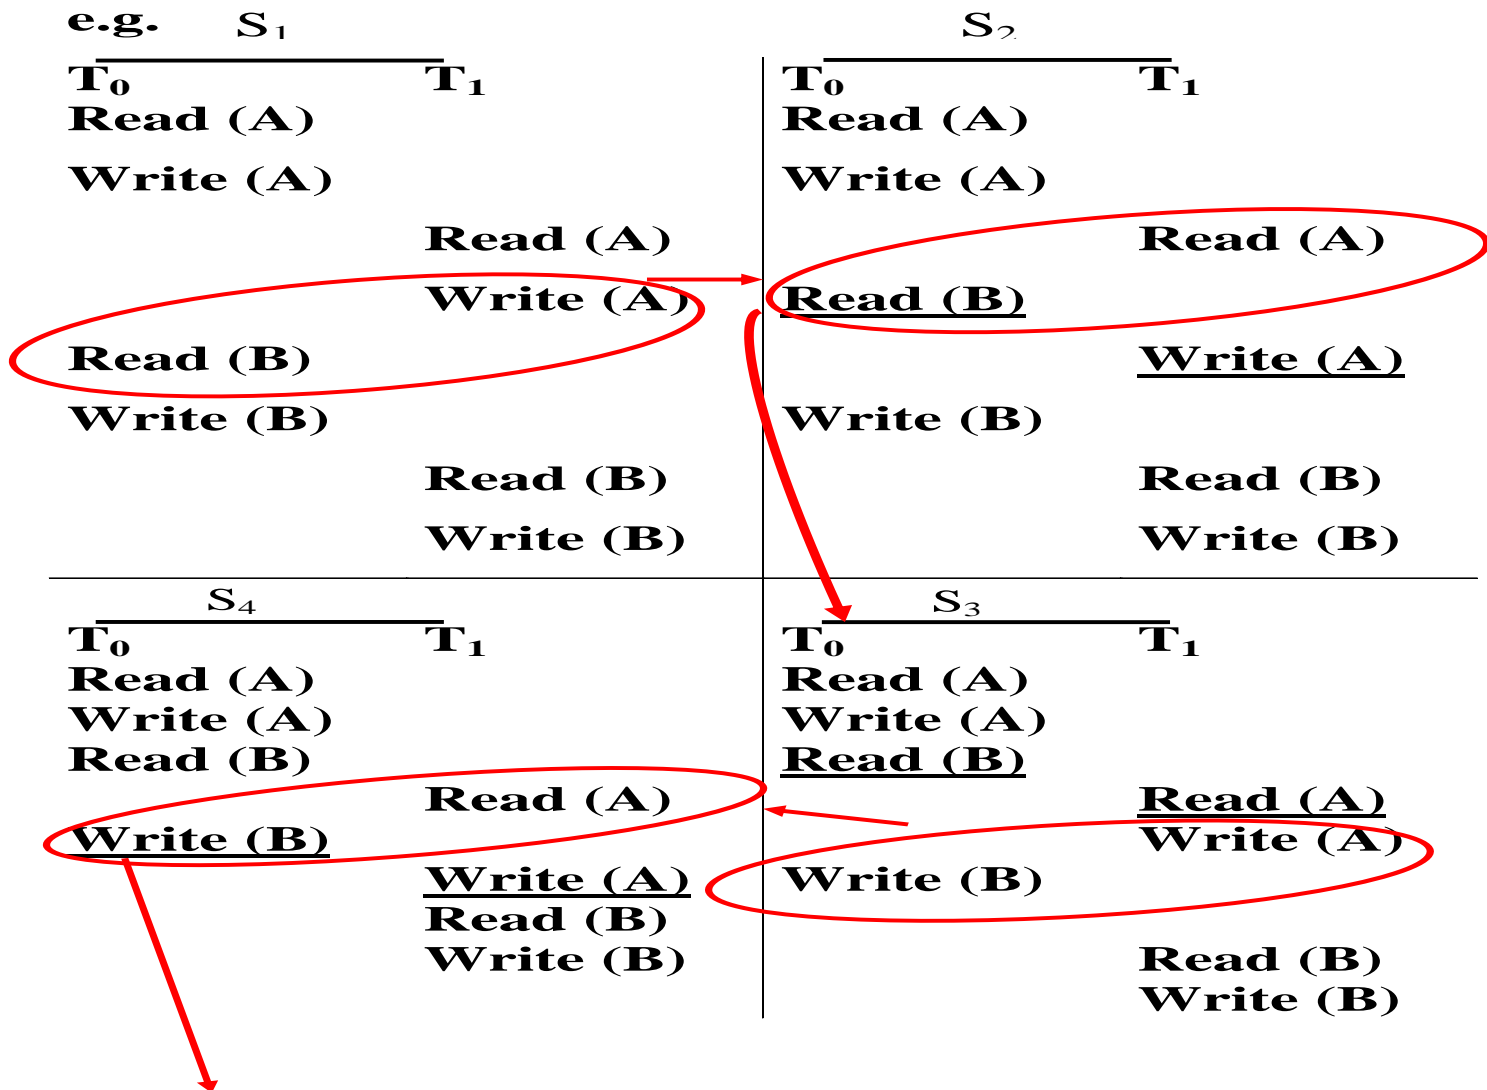
## SCHEDULES (9)

e.g.




$$S \equiv S'$$

# SCHEDULES (10)



# SCHEDULES (11)



**Serial Schedule**

**$S_1 \equiv S_2 \equiv S_3 \equiv S_4 \equiv S_5$**

**$\Rightarrow S_1, S_2, S_3, S_4 : \text{Serializable Schedules}$**

# CONCURRENCY CONTROL TECHNIQUES

- To ensure serializability of transaction schedules
- Three major types:
  - Lock-based techniques
  - Timestamp-based techniques
  - Validation-based techniques



# CONCURRENCY CONTROL TECHNIQUES (2)

## a) Lock-based Techniques:

- Allow a trans to access a data item only if it is currently **holding a lock** on that item.
  - Example: 2-phase locking, graph-based locking
- Locks: 2 types
  - Shared Lock: S
    - If trans T holds an S lock on data item Q => T **can read** Q but **cannot write** Q
  - Exclusive Lock: X
    - If trans T holds an X lock on data item Q => T **can read** Q and **can write** Q

# CONCURRENCY CONTROL TECHNIQUES (3)

## Lock Compatibility Matrix

		Locks requested by $T_j$	
		S	X
Locks held by $T_i$	S	TRUE	False
	X	False	False

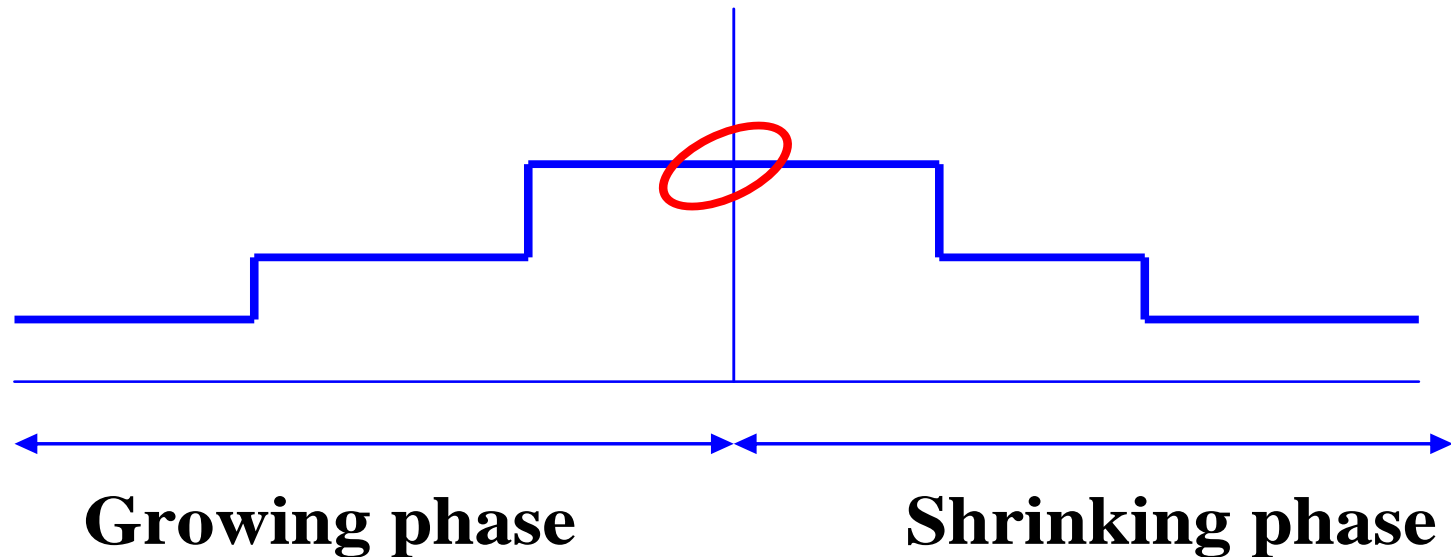
# CONCURRENCY CONTROL TECHNIQUES (4)

- A transaction requests an **S lock on Q** by executing **lock\_S (Q)**
- A transaction requests an **X lock on Q** by executing **lock\_X (Q)**
- A transaction **unlocks** a data item **Q** by executing **unlock (Q)**
- To access a data item **Q**, trans  $T_i$  must:
  - request a lock on **Q**
    - if **Q** is currently locked by another trans in an incompatible mode =>  $T_i$  must **wait**.

# CONCURRENCY CONTROL TECHNIQUES (5)

- Two-phase locking Technique:
  - Ensures serializability
  - Each trans issues **lock** and **unlock requests** in 2 phases
    - **Growing phase:** A trans may obtain locks but may not release any lock
    - **Shrinking phase:** A trans may release locks but may not obtain any new locks

# CONCURRENCY CONTROL TECHNIQUES (6)



**This technique: does not ensure freedom from deadlock**

# CONCURRENCY CONTROL TECHNIQUES (7)

- Example:



# CONCURRENCY CONTROL TECHNIQUES (8)

## b) Timestamp-based techniques:

- Determine the serializability order by selecting an ordering among transactions using timestamps in advance.
- Example:
  - Timestamp-ordering scheme
  - Thomas' write rule



# CONCURRENCY CONTROL TECHNIQUES (9)

- Timestamp:
  - Associate a unique timestamp **TS** with each trans  $T_i$  **TS( $T_i$ )** when  $T_i$  enters the system
  - If trans  $T_j$  enters the system after  $T_i$ , then:
$$TS(T_j) > TS(T_i)$$
  - Associate with each data item  $X$  two timestamp values:
    - **W\_TS ( $X$ ): Write Timestamp of  $X$ :** largest timestamp among all timestamps of trans that have successfully executed Write( $X$ ).
    - **R\_TS ( $X$ ):** similar for **Read Timestamp of  $X$**

# CONCURRENCY CONTROL TECHNIQUES (10)

- **Timestamp-Ordering Technique:**
  - Orders trans based on their timestamps
  - A serializable schedule: equivalent to a serial schedule that corresponds to the order of trans timestamps
  - Ensures that any conflicting READ and WRITE operations are executed in the timestamp order as follows:

# CONCURRENCY CONTROL TECHNIQUES (11)

- When  $T_i$  issues READ (X), check:
  - If  $TS(T_i) < W\_TS(X)$ :
  - If  $TS(T_i) \geq W\_TS(X)$ :

# CONCURRENCY CONTROL TECHNIQUES (12)

- When  $T_i$  issues WRITE (X), check:
  - If  $TS(T_i) < R\_TS(X)$ :
  - If  $TS(T_i) < W\_TS(X)$ :
  - Otherwise:
- Note: in the Timestamp-ordering scheme:
  - There is no deadlock
  - Cascading rollback is possible -> Why?

# CONCURRENCY CONTROL TECHNIQUES (13)

## c) Validation-Based Techniques (Optimistic):

- No conflict checking is done during trans execution
- Each trans goes through 2 or 3 phases:
  - 1-READ PHASE:
- At the end of trans execution:
  - 2-VALIDATION PHASE:
  - 3-WRITE PHASE:
- Phase 1+2: for what trans?
- Phases 1+2+3: for what trans?

**END OF TOPIC**  
**“Concurrency Control”**